

Presentation Unsupervised Learning project

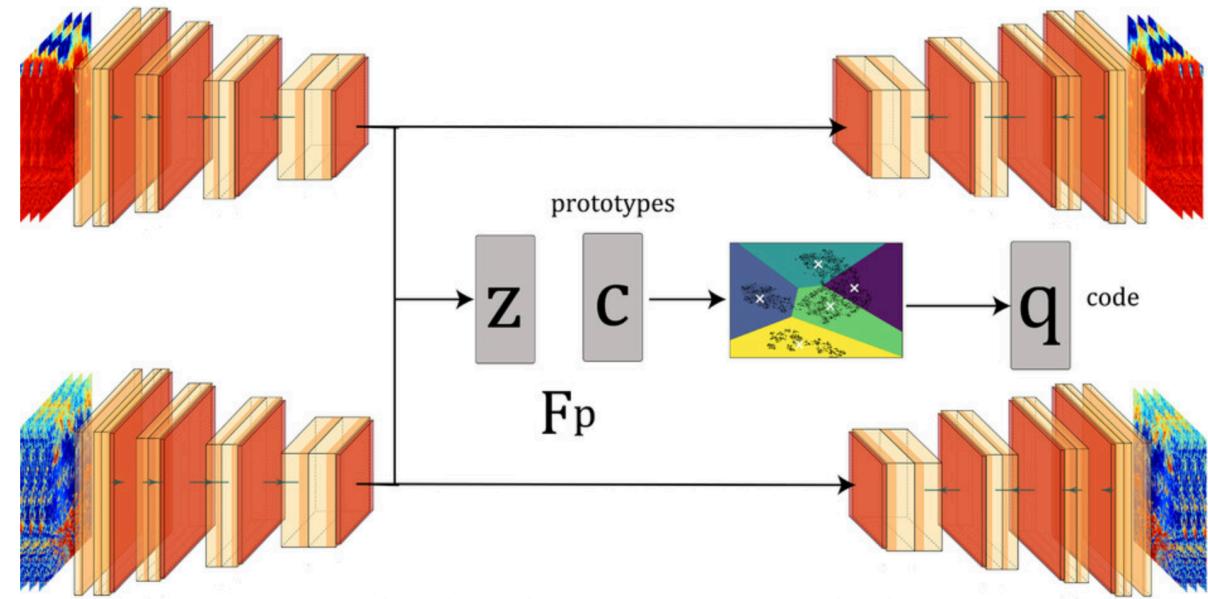
Marco Zampar - SM3800032

January 13, 2026

Clustering Auto-Encoder

This project is aimed at implementing the Clustering Autoencoder (CAE), a method proposed in: **Identifying Climate Patterns Using Clustering Autoencoder Techniques [1]**

The idea is to add to the Convolutional Autoencoder a **single-layer perceptron** F_p to map the latent representation z into \mathbb{R}^n (the space where the **cluster prototypes** c are defined) and to return a probability distribution over the possible clusters proportional to the scalar product $F_p(z) \cdot c$.



Training phase

Given a data point x , the *CAE* encodes the data in a **latent space** through an **encoder**, this latent representation is flattened into a vector $z \in \mathbb{R}^m$.

The latent representation z is passed through the **decoder**, the reconstructed \hat{x} is obtained and the **reconstruction loss** $RMSE(x, \hat{x})$ is computed.

The latent representation z is also passed through a **single-layer perceptron** with a **ReLU** activation function to map it into \mathbb{R}^n . Given a set of k (trainable) **cluster prototypes**, the probability distribution of a point x of belonging to a given cluster is computed passing the scalar product $F_p(z) \cdot c$ through a **softmax layer**, with temperature parameter τ that regulates the probability of the most probable cluster, (the lower τ , the higher the probability of the most probable cluster).

$$p(k_i) = \frac{\exp\left(\frac{F_p(z) \cdot c_i}{\tau}\right)}{\sum_{j=1}^k \exp\left(\frac{F_p(z) \cdot c_j}{\tau}\right)}$$

Target cluster assignments are obtained with the iterative **Sinkhorn-Knopp algorithm** [2][3]: given a batch of B samples, it assigns a probability distribution over clusters to each data point in an efficient, non trivial and balanced way, with a smoothing parameter ϵ .

Let us define: $Z \in \mathbb{R}^{B \times n}$, the latent representation of the points of the batch, passed through F_p ; $C \in \mathbb{R}^{n \times k}$, the cluster prototypes; $Q \in \mathbb{R}^{k \times B}$, the cluster probability distribution for each data point in the batch.

The matrix $Q = \text{Diag}(\mathbf{u}) \exp\left(\frac{C^\top Z}{\epsilon}\right) \text{Diag}(\mathbf{v})$, where u and v are computed by renormalization in \mathbb{R}^k and \mathbb{R}^B through the Sinkhorn–Knopp algorithm (iterative row-normalisation and column-normalisation), is the solution to the problem:

$$\max_{Q \in \mathcal{Q}} \text{Tr}(Q^\top C^\top Z) + \varepsilon H(Q)$$

where H is the entropy function and the possible matrices $Q \in \mathcal{Q}$ satisfy the conditions:
 $Q\mathbf{1}_B = \frac{1}{K}\mathbf{1}_K$, $Q^\top \mathbf{1}_K = \frac{1}{B}\mathbf{1}_B$, allowing a balanced cluster assignment over the batch, with no collapse of assignments to a single cluster.

```
def sinkhorn(scores, epsilon=0.05, sinkhorn_iterations=3):
    # scores: (B, feature_dim)
    with torch.no_grad():

        Q = torch.exp(scores / epsilon).t()  # K x B
        B = Q.shape[1]
        K = Q.shape[0]

        Q /= torch.sum(Q)

        for _ in range(sinkhorn_iterations):
            # row normalization
            Q /= torch.sum(Q, dim=1, keepdim=True)
            Q /= K

            # column normalization
            Q /= torch.sum(Q, dim=0, keepdim=True)
            Q /= B

        Q *= B
    return Q.t()  # back to (B, K)
```

The **clustering loss** is computed via **cross-entropy** between the target assingment distributions and predicted assingment distributions.

The final loss is the **weighted sum** of the reconstruction and clustering loss terms. When we optimise it by Gradient descent, we optimse both the encoder-decoder architectures, the cluster prototypes and the parameters of the single-layer perceptron F_p .

The encoder and decoder architectures, since we are working with 3D data (lat x lon x time), are **Convolutional Neural Networks**. We kept a similar structure of the CNN encoder and decoder wrt the orignal paper.

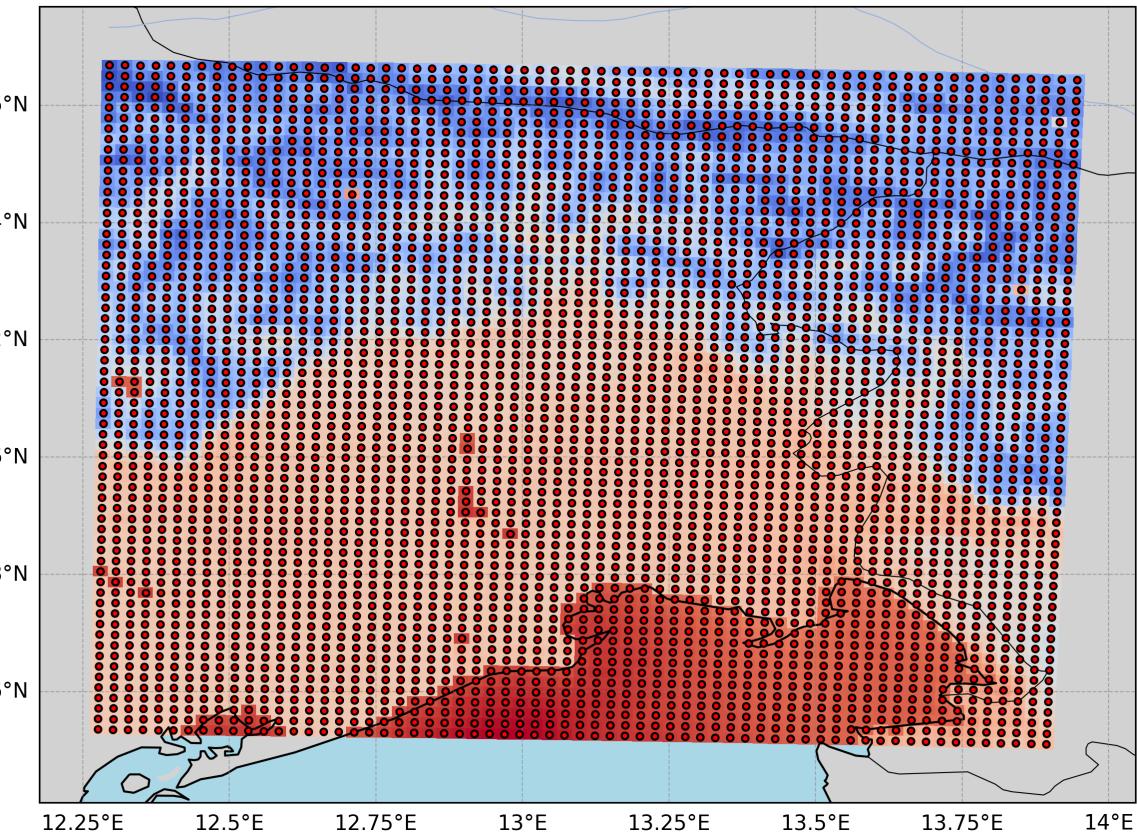
Experiment:

Are there any climatological regional patterns in the **Friuli Venezia Giulia** region that can be found through a clustering technique?

Working with a reanalysis of 2km spatial resolution ($68 \times 66 = 4488$ gridpoints) and hourly temporal resolution, we analyse more than 20 years (2001-2023) of **precipitation** (mm/h), working with **monthly means**.

The period (2001-2017) is taken as training set and the period (2018-2023) as test set.

When talking about **clusters**, we refer to a partition of the gridpoints based on similar **climatological patterns**.



Data and Training:

After rescaling the data (min-max scaling on a monthly basis) and extracting patches of shape 8x8 (64 gridpoints) moving the patch with a **stride** of 2, the monthly cumulatives of each year are stacked one on the other, to get a 3D-tensor of shape **8x8x12**.

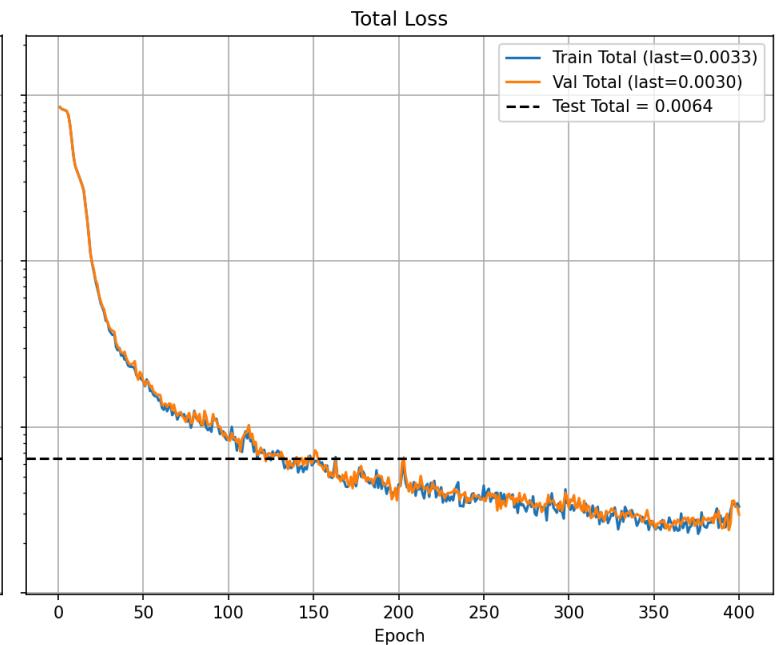
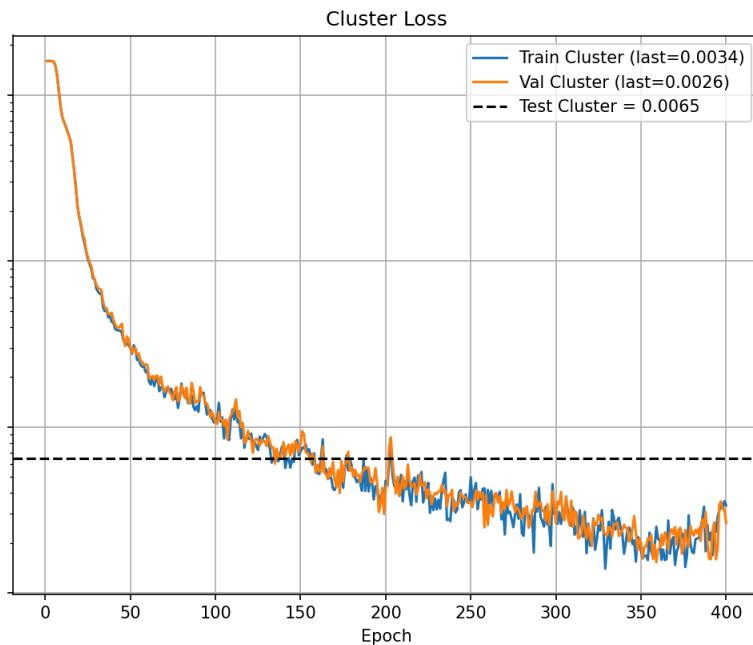
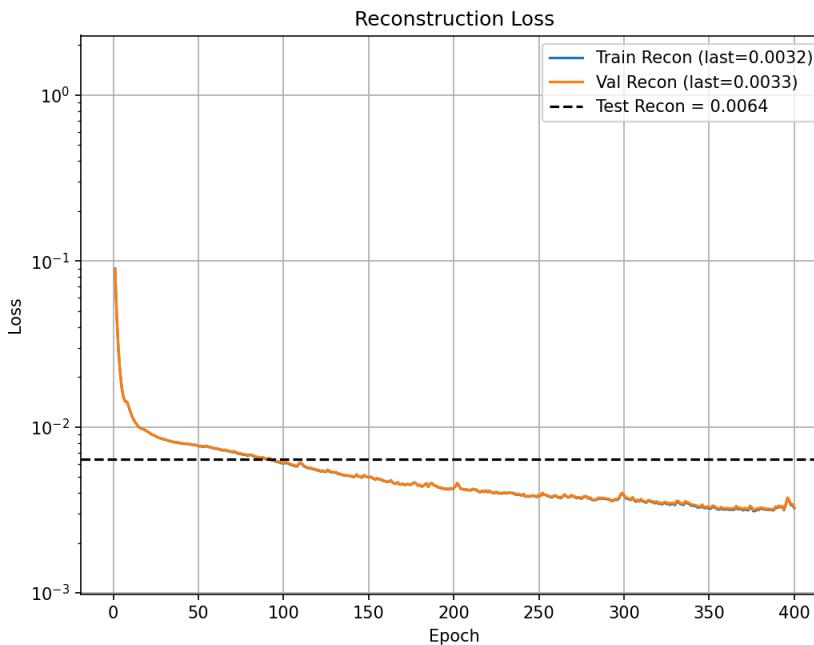
We get a training set of **15810** patches and a test set of **5580** patches.

Each grid cell is assigned to the **most frequent cluster** of the patches that contain that grid cell.

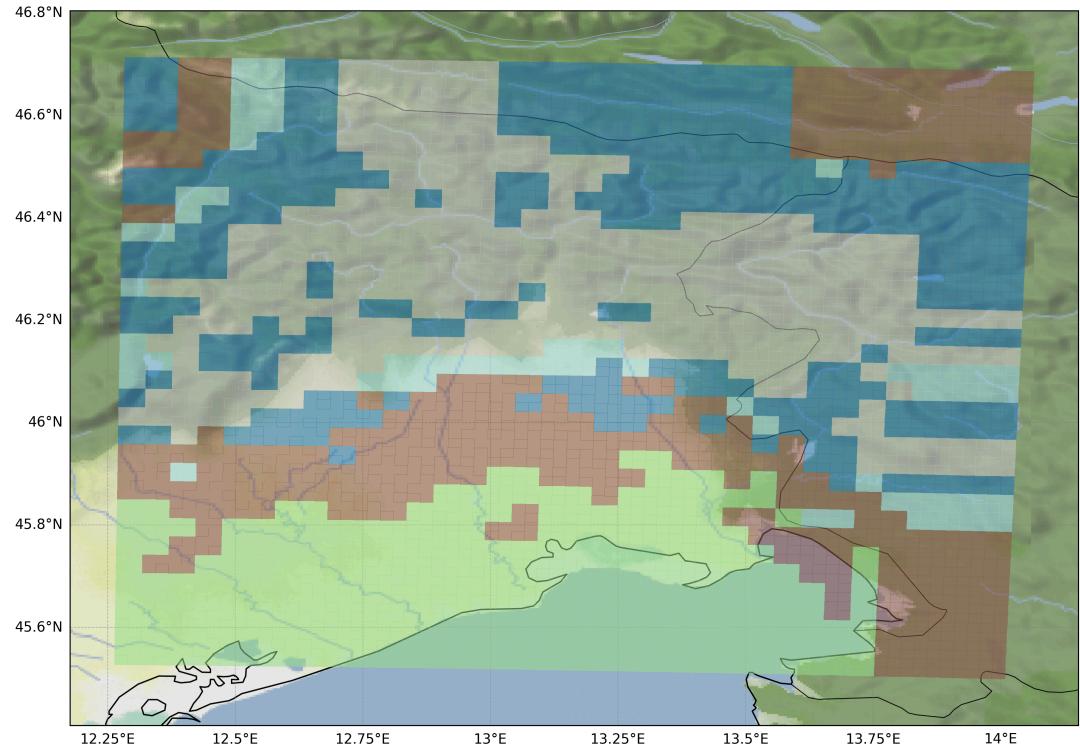
After some tuning of the parameters, we found a nice configuration (different from the one of the paper):

- $\alpha = 0.5$
- $\tau = 0.05$
- $\epsilon = 0.09$
- learning rate = 10^{-4}
- batch size = 512
- latent dimension = 512
- prototypes dimension = 512

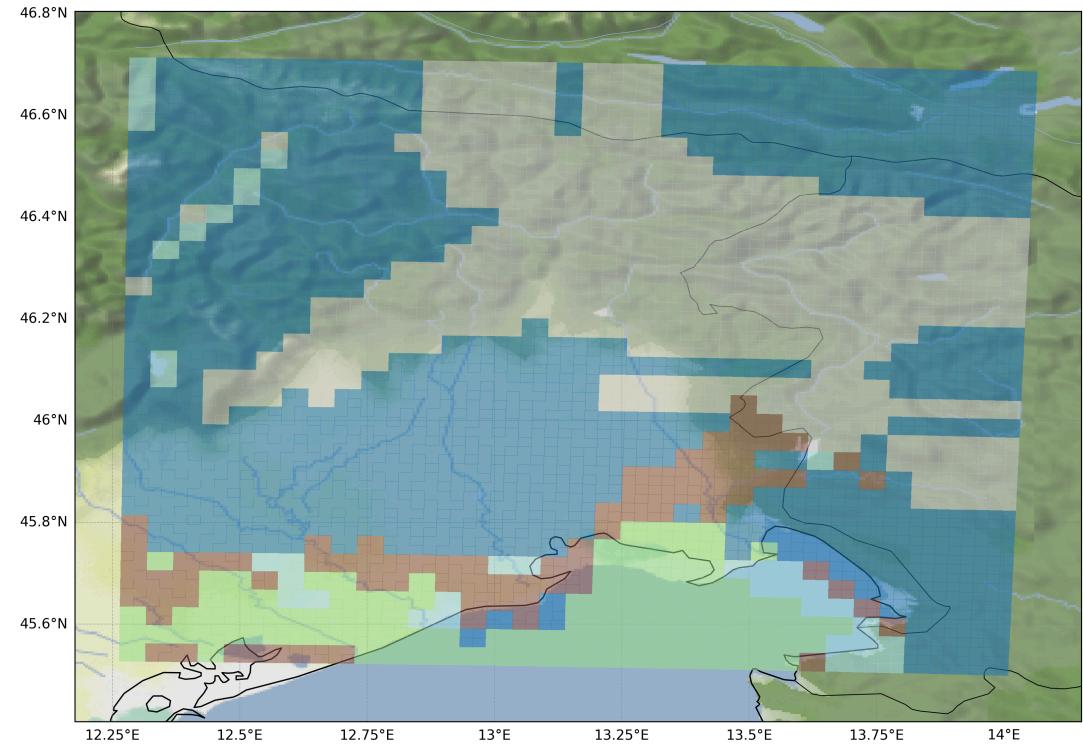
Training History



Training clustering



Test clustering



Comparison with other methods:

A different method was tested: **k-medoids**. This was chosen because, despite its simplicity, it allows to use a generic distance matrix.

Many number of clusters were tested, and since the results can be easily visually inspected, the choice of the number of clusters was done looking at the cluster plots.

2 distances were used: the **Pearson correlation coefficient** between each pair of monthly time series ($d_{ij} = 1 - \rho$) and the **Kolmogorov-Smirnov distance** between the **Empirical Cumulative Distribution Functions** of each pair of samples ($d_{ij} = D_{KS}(F_i, F_j)$).

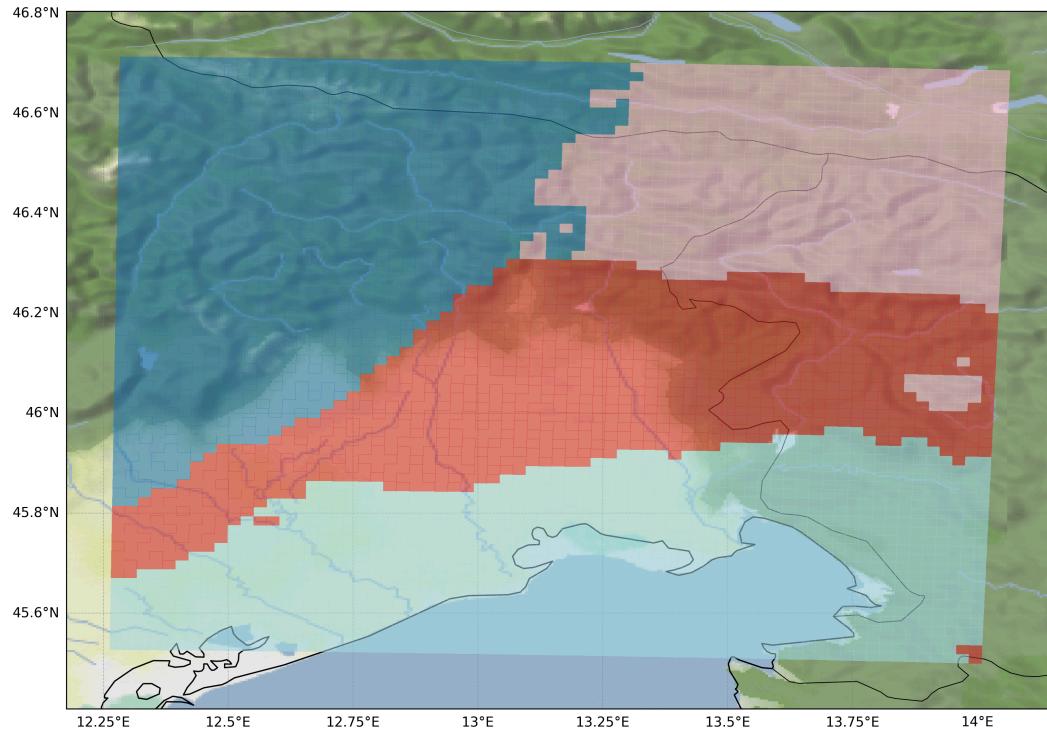
The **KS distance** between two cumulative distribution functions (CDFs) ($F(x)$) and ($G(x)$) is defined as:

$$D_{KS} = \sup_x |F(x) - G(x)|$$

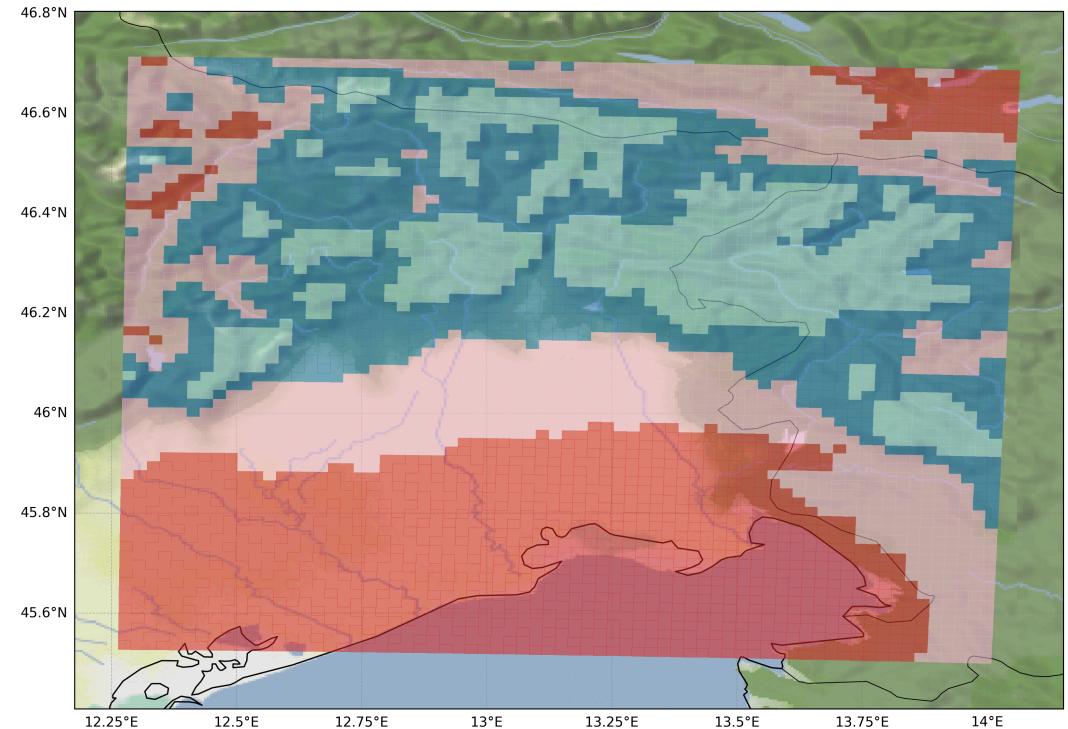
where \sup_x denotes the maximum absolute difference between the two CDFs.

K-medoids clustering

Correlation-based



KS-based



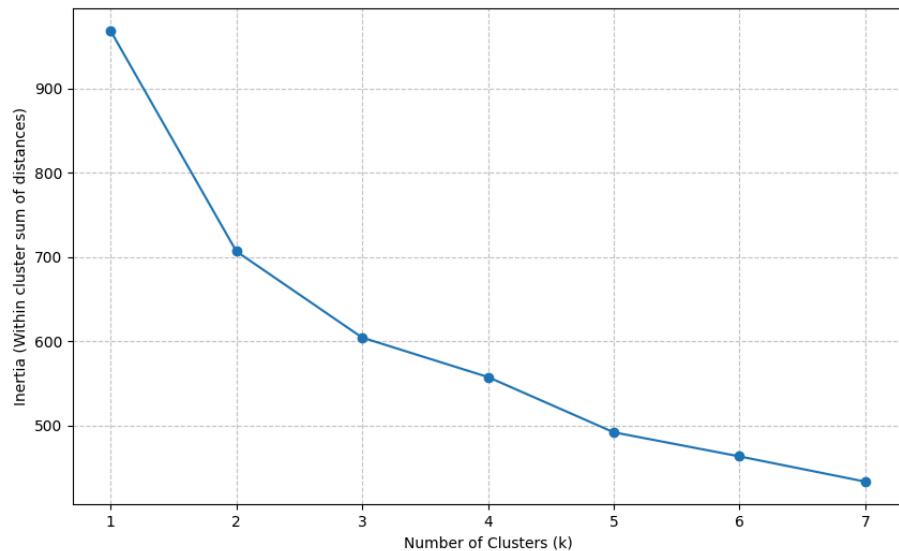
In the first case we find homogenous-in-space clusters, probably due to the fact that precipitation shows high **spatial autocorrelation**.

In the second case, since KS distance doesn't take into account the temporal correlation of the data, but only the distance between its **ECDFs**, we can see that the results seem to have a physical meaning and are probably driven by the orography.

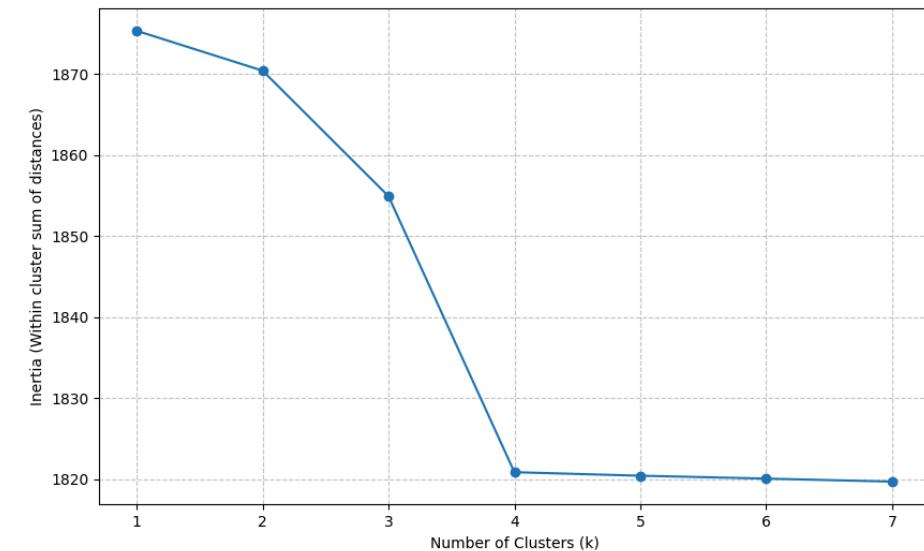
The number of clusters is chosen to be **4**. This was chosen looking at the KS distance-based cluster plots, promoting interpretability and physical meaning of the clusters, and the Scree test.

Scree test

Correlation-based k-medoids



KS-based k-medoids



Pros and cons of the CAE

Strengths

- Leverages more information: k-medoids condenses all information into a single distance matrix, while CAE uses information coming from different years.
- Takes spatial autocorrelation into account with 8x8 pixel patches, but pools together different years of the same patch, allowing robust yet non-dominant spatial patterns, combining KS distance and Pearson correlation approaches.
- Scales well to large datasets.
- Fuzzy clustering can be obtained by combining information of the two methods.

Weaknesses

- Requires a GPU.
- Overfitting: model may not generalize well.
- Unreliable results: training sometimes fails to converge.
- Initialization dependence: different runs give different results.
- Parameter tuning (tau and epsilon) is non-obvious: using paper parameters didn't work.
- Number of clusters is a

Conclusions and further experiments

In this brief presentation, we outlined the main characteristics of the Clustering Auto-Encoder, its pros and cons and the differences wrt the outputs of other clustering techniques.

In the future, we can work on these aspects:

- make the model more reliable and stable, facing the overfitting and initialisations dependence issues;
- apply this technique to a bigger dataset in terms of gridpoints and years or increase the same dataset using a different temporal aggregation strategy, like decades instead of months;
- apply this technique to other meteorological variables, like temperature and relative humidity;
- apply this technique to the future scenario projections CMIP6 Eurocordex.

Thank you for your attention!

References

1. Kurihana et al., 2024, **Identifying Climate Patterns Using Clustering Autoencoder Techniques**: <https://journals.ametsoc.org/view/journals/aies/3/3/AIES-D-23-0035.1.xml>
2. Caron et al., 2020, **Unsupervised learning of visual features by contrasting cluster assignments**: <https://arxiv.org/pdf/2006.09882>
3. Cuturi, 2013, **Sinkhorn distances: Lightspeed computation of optimal transport**: <https://arxiv.org/abs/1306.0895>